2026-05-04 21:10:56 You Asked

Gemini 3.1 Pro

This is a fascinating graph. You have successfully trained the model, but you caught a "rogue" batch right in the middle of the run.

Here is exactly what this dashboard is telling us, and why that massive purple square at Step 700 happened:

1. The Core Training is a Success

Ignore the purple lines for a second and look purely at your orange **Train loss** and blue **Train mean token accuracy**.

- They are beautiful. The training loss settled right into that **1.3** pocket, and the accuracy is consistently hovering just under **70%**.
- Because the orange line never spikes, we know the model successfully learned the mathematical rules and formatting of your cleaned data. It didn't experience "catastrophic forgetting" or model collapse. It knows exactly what it is doing.

2. The Step 700 Anomaly (The Poison Pill)

Look at the purple **Valid loss** line. It's cruising along normally around 2.0, and then at exactly Step 700, it violently spikes up to **4.5** (and the validation accuracy plummets down to **40%**). At Step 800, it drops right back down to normal.

What does this mean? During training, Azure doesn't test the model on the *entire* validation set every 100 steps. It pulls a random "batch" from your validation file. At Step 700, the model grabbed a validation batch, looked at the `USER` prompt, tried to predict the `ASSISTANT` notes, and completely failed. It was penalized heavily.

Because the train loss didn't flinch, the model itself didn't break. **It just hit a dirty file.**

3. Why did it happen if you cleaned the validation data?

Your cleaning script likely did a great job on 95% of the files, but an outlier slipped through. There is almost certainly a file (or a few files) in your `validation.jsonl` where the LLM cleaner failed to strip out a massive chunk of human gossip, or the Trial Balance math is so bizarre that it doesn't align with the notes at all. When the model was tested against that specific file at Step 700, it got whiplash.

The Final Verdict: Is the model usable?

Yes, absolutely.

Look at the magenta line (`Full valid loss`). This is Azure testing the model against the *entire* validation set at major checkpoints, rather than just random batches. That line stays completely stable. That tells us the model handles the vast majority of your validation data perfectly.

You do not need to retrain. You have a highly competent, mathematically grounded AI agent sitting in your deployments right now.

Go to your deployments tab, load up a fresh Xero Trial Balance that this model has never seen, and ask it to write the working papers. I think you are going to be thrilled with the output!